

**AlwaysON contained availability  
groups, What?**

**DATA**  
SATURDAYS



## Sponsors



With the support of:




# About me

- Massimiliano Buschi [mbuschi@lucient.com](mailto:mbuschi@lucient.com)
- Data Platform Architect @ Lucient Italia One4 Group
- 15 + years of work with SQL Server
- SQL & Data Saturday lurker
- [DBA Stack Exchange Community Member](#)
- Speaker here for the first time!



# Always On availability groups: a high-availability and disaster-recovery solution

- **High-availability** and **disaster-recovery solution (AG)**
- Introduced in SQL Server 2012
- Maximizes the **availability** of **a set of user databases**
- Supports a set of read-write primary databases and one to eight sets of corresponding secondary databases
- AlwaysOn AG is **not** a Backup Solution (like true HA&DR solutions)
- Requires a Windows Server Failover Cluster (**WSFC**)
  - a cluster role is created for every availability group that you create



With the support of:

## Always On availability groups: benefits

- Supports **up to nine** availability replicas
- 2 [replication modes](#):
  - **Asynchronous**-commit
  - **Synchronous**-commit (5 replicas since 2019, before only 3)
- Several forms of availability-group [failovers](#):
  - **Automatic** failover
  - **Planned** manual
  - **Forced** manual
- You can configure a given availability replica to support the following active-**secondary capabilities**:
  - **Read-only** connections
  - [Backup operations](#)
- Supports an availability group [listener](#) for each availability group
- Supports [automatic page repair](#) for protection against page corruption
- Supports **encryption** and **compression**



# Always On availability groups: can we have them in SQL Server Standard Edition?

- Sadly, NO or better you can have a simpler version called [Basic Availability group \(BAG\)](#)
- These are the **differences**:
  - **Only two replicas**
    - The primary and the secondary replica
  - You can **use only the primary replica**
  - Support for one availability database
    - the 'group' can be composed by only **a single database**

# Always On availability groups: how to manage them?

- We have a set of tools to simplify deployment and management of availability groups, including:
  - [Transact-SQL DDL statements](#)
  - **SQL Server Management Studio** (for SQL server 2022 use [SSMS 19.0.1](#))
    - [The New Availability Group Wizard](#)
    - [The Add Database to Availability Group Wizard](#)
    - [The Add Replica to Availability Group Wizard](#)
    - [The Fail Over Availability Group Wizard](#)
  - [The Always On Dashboard](#)
  - [PowerShell](#) and [DBATools cmdlets](#)




## Always On availability groups: client connectivity with availability group listener

- An availability group **listener** is a **server name** (and an IP number) to which **clients can connect** in order to access a database in a primary or secondary replica of an Always On availability group.
- Availability group listeners **direct** incoming connections **to the primary** replica or **to a read-only secondary** replica.
- The listener provides **fast application failover** after an availability group fails over.
- They rely on **failover cluster services role switch** to work



- AG Setup
- Failover
- Connectivity to listener

## Alwayson AG: before sql 2022



- We have some issues
  - A few types of objects are outside AG replication as they are instance level objects
  - Login
  - Jobs
  - Linked server
  - Anything in the master database
  - Anything in the msdb database
    - Alerts
    - Operators

- A **contained availability group (CAG)** is an Always On availability group that supports:
  - **managing metadata** objects (users, logins, permissions, SQL Agent jobs etc.) **at the availability group level in addition to the instance level.**
  - specialized **contained system databases** within the availability group.

```
SQL Copy  
  
<with_option_spec>::=  
CONTAINED |  
REUSE_SYSTEM_DATABASES
```

Specify availability group options

Availability group name:

Cluster type: Windows Server Failover Cluster ▼

☐ Database Level Health Detection

☐ Per Database DTC Support

☒ Contained

☒ Reuse System Databases

During CAG creation sysadmin logins at instance level are copied to the contained ag master database

- And other logins?
  - **They are not copied**
- What will happen next ?
  - **NOTHING will be copied automatically.**

Differences between connecting to the instance and connecting to the contained availability group:

- When connected to contained AG, **users will only see databases in the contained AG, plus tempdb.** 👍
- At instance level, contained AG master and msdb names will be **[contained AG]\_master**, and **[contained AG]\_msdb**. Inside contained AG, their names are master and msdb. 👍
- **Database ID for contained AG master is 1** from inside contained AG, but something else when connected to the instance. 👍
- While **users will not see databases outside of the contained AG** in sys.databases when connected in a contained AG connection, **they will be able to access those databases by three part name** or through the use command. 👍
- Server configuration through sp\_configure can be read from contained AG connection but **can only be written from instance level.** 👍
- From contained AG connections, **sysadmin is able to perform instance level operations**, such as shutting down SQL Server. 👍
- Most DB level, end point level, or AG level operations **can only be performed from instance connections**, not contained AG connections. 👍



- Contained AG Setup
- Failover
- Connectivity to listener
  - Differences
    - Logins
    - Jobs
    - Linked server

## Good things:

- **You don't need to keep login/jobs in sync** between nodes and manage job running on secondaries
- **Filtered database list in SSMS** gives and idea of AG boundaries.

## Things that need improvements:

- TDE: you have to **manually install the DMK** into the contained master databases

## Things **not** currently **supported**:

- SQL Server Replication
- Distributed AG
- Log shipping where the target database is in the contained availability group

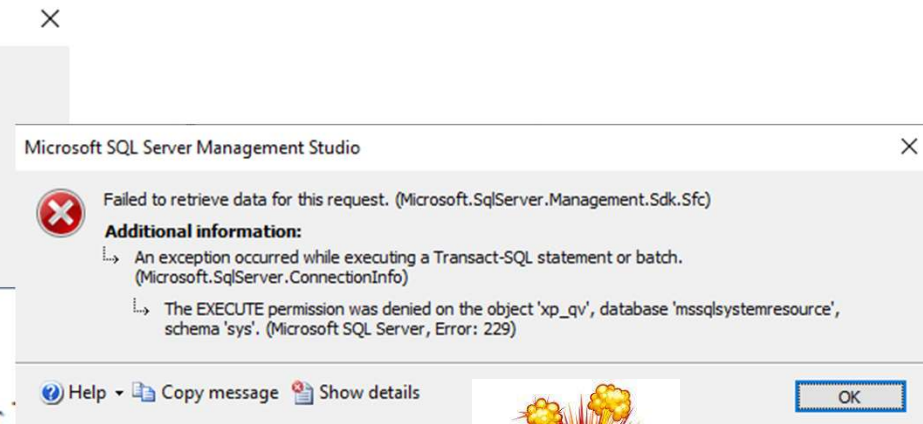
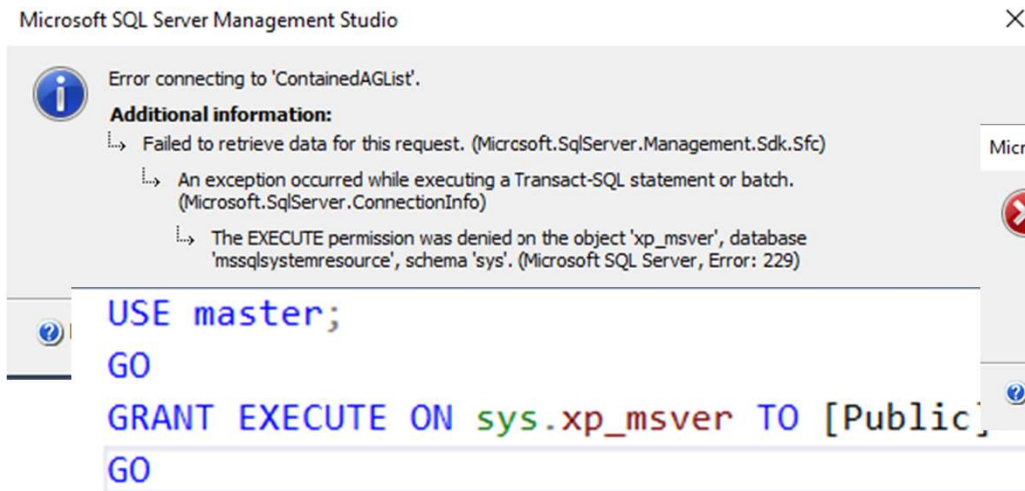
## Things to pay attention to:

- Define a **contained objects naming convention**; something like **agname\_jobname, agname\_loginname** etc.. as the contained system databases.
- Administer contained objects connecting to the listener only, **anything else connecting only to instances**
- **Do not rely** on the possibility of **querying databases outside the ag**, especially for data. If you need them, move them to the group or merge the two groups.
- No instance level object are synchronized between contained ag 'world' and instance 'world' after creation. **You have to do it manually.**



## Last minute errors!

- The execute permission was denied on the object xp\_msver. Database mssqlsystemresource, schema sys. (Microsoft SQL Server, Error 229)



**Thank You All**

**DATA**  
**SATURDAYS**





May the forced  
plan be with you

**DATA**  
SATURDAYS

